

LLMs | RAG | AI Agents

Step-by-step tutorials



Awesome LLM Apps

100+ open-source AI agents, agent skills, and RAG apps. Hand-built, tested end-to-end, Apache-2.0.

Clone it, ship it, sell it - 100% free and open-source

Works with Claude, Gemini, GPT, DeepSeek, Llama, Qwen and other open-source models.

[Step-by-step tutorials on Unwind AI](#) · [Quick start](#) · [Browse all templates](#)




GITHUB TRENDING

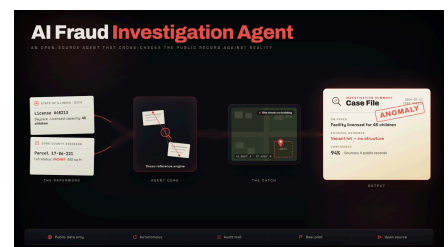
#1 Repository Of The Day



Project Graveyard



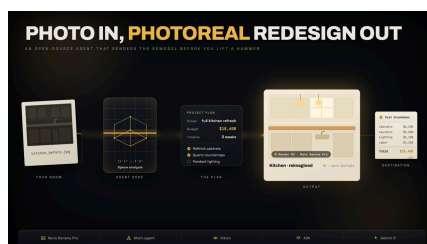
Insurance Claim Live Agent Team



AI Fraud Investigation Agent



Self-Improving Agent Skills



AI Home Renovation Agent



Always-on HN Briefing Agent



Run one now

Give your coding agent a new skill in 10 seconds:

```
npx skills add https://github.com/Shubhamsaboo/awesome-llm-apps/tree/main/agent_skills/project-{"
```

Then ask it: *"why do I never finish my side projects?"*

Or clone and run any agent in 30 seconds:

```
git clone https://github.com/Shubhamsaboo/awesome-llm-apps.git
cd awesome-llm-apps/starter_ai_agents/ai_travel_agent
pip install -r requirements.txt
streamlit run travel_agent.py
```



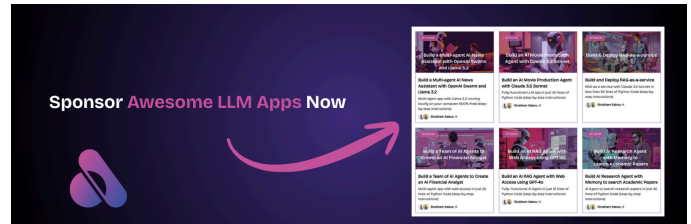
New templates drop weekly. [Get them in your inbox on Unwind AI](#).



Thanks to our sponsors



Vorflux



Become a Sponsor



Browse all templates



Agent Skills

Give your coding agent new abilities. One command to install, plain English to use. Every skill ships real code and passes a security + eval CI gate. Works with Claude Code, Codex, Cursor, and other coding agents. [Browse all skills](#) →

-  [Project Graveyard](#) - Finds every side project you abandoned, tells you why each one died, and helps you finish the one worth going back to

- 🕸 [Scope Creep Detector](#) - Checks whether a diff grew beyond its stated intent and recommends what to keep, split, or justify
- 🏺 [Commit Archaeologist](#) - Reconstructs why a file or code region exists from its introducing commit, later edits, co-changes, and intent clues
- 🧰 [Dependency Doctor](#) - Checks a dependency manifest for standard-library pins, obsolete backports, unpinned entries, duplicate constraints, and yanked releases
- 🗣 [Advisor Orchestrator Worker](#) - Meta Loop with Claude Fable 5 as advisor, GPT-5.6 as orchestrator, and Gemini 3.7 Flash as worker
- ∞ [Self-Improving Agent Skills](#) - Automatically optimize agent skills using Gemini and ADK

Starter AI Agents

Single-file agents that run with just an API key - a great place to start.

- 🎙 [AI Blog to Podcast Agent](#) - Turn any blog URL into a narrated podcast episode
- ❤ [AI Breakup Recovery Agent](#) - An agent team that talks you through the post-breakup spiral
- 📊 [AI Data Analysis Agent](#) - Ask questions of any CSV or Excel file in plain English
- 🧠 [AI Medical Imaging Agent](#) - Diagnostic analysis of X-rays and scans with Gemini
- 🤪 [AI Meme Generator Agent \(Browser\)](#) - Makes memes by driving a real browser, not an image API
- 🎵 [AI Music Generator Agent](#) - Prompt in, MP3 track out
- ✈ [AI Travel Agent \(Local & Cloud\)](#) - Personalized day-by-day travel itineraries
- 💰 [AI x402 Paying Agent](#) - An agent with a wallet that pays per-call for the data it needs — no API keys
- ✨ [Gemini Multimodal Agent](#) - Video analysis plus web search in one agent
- 🔄 [Mixture of Agents](#) - Multiple LLMs answer, one aggregates the best response
- 📈 [xAI Finance Agent](#) - Real-time stock analysis powered by Grok
- 🔍 [OpenAI Research Agent](#) - Multi-agent topic research with the OpenAI Agents SDK
- ⚙ [Web Scraping AI Agent](#) - Describe what to extract and the agent scrapes it

Advanced AI Agents

Production-style agents with tools, memory, and multi-step reasoning.

- 🏠 🌟 [AI Home Renovation Agent with Nano Banana Pro](#) - Photos of your space in, renovation plan and photorealistic renders out
- 🗣 [DevPulse AI - Multi-Agent Signal Intelligence](#) - Aggregates and scores technical signals into a daily intelligence digest

-  [AI Deep Research Agent](#) - Comprehensive web research with the OpenAI Agents SDK and Firecrawl
-  [AI VC Due Diligence Agent Team](#) - Multi-agent startup investment analysis with Gemini 3
-  [AI Research Planner & Executor \(Google Interactions API\)](#) - Multi-phase research with stateful conversations and auto-generated infographics
-  [AI Consultant Agent](#) - Market analysis and strategy recommendations with live web research
-  [AI System Architect Agent](#) - Architecture reviews using DeepSeek R1 reasoning plus Claude
-  [AI Financial Coach Agent](#) - Personalized budget, debt, and savings analysis
-  [AI Movie Production Agent](#) - Script drafts and casting ideas from a one-line movie concept
-  [AI Investment Agent](#) - Stock comparison reports built on Yahoo Finance data
-  [Earnings Call Analyst Agent](#) - Turns YouTube earnings calls into a playback-synced analyst workspace
-  [AI Health & Fitness Agent](#) - Tailored diet and workout plans from your goals
-  [AI Product Launch Intelligence Agent](#) - Go-to-market intelligence on competitor launches
-  [AI Fraud Investigation Agent](#) - Cross-references public records to flag facilities that don't add up
-  [AI Journalist Agent](#) - Researches, writes, and edits articles on any topic
-  [AI Mental Wellbeing Agent](#) - A coordinated agent team for mental health support plans
-  [AI Meeting Agent](#) - Context, industry insights, and strategy briefs before you walk in
-  [AI Self-Evolving Agent](#) - Agents that rewrite their own workflows with EvoAgentX
-  [AI Sales Intelligence Agent Team](#) - Generates competitive sales battle cards in real time
-  [AI Social Media News and Podcast Agent](#) - Curates your trusted sources into briefs and generated podcasts
-  [Openwork - Open Browser Automation Agent](#) ^{external} - Open-source agent that operates a real browser
-  [Trust-Gated Multi-Agent Research Team](#) - Every agent verified, every action in a hash-chained audit trail

Always-on Agents

Background agents that run on schedules or events, monitor changing context, decide what needs attention, and proactively deliver updates, artifacts, or actions.

-  [Always-on Hacker News Briefing Agent](#) - A scheduled scout that ships a ranked daily brief to Slack or email
-  [Release Radar Agent](#) - Watches dependency releases and briefs you on breaking, deprecated, security, and major-version changes

Multi-agent Teams

Multiple agents collaborating to accomplish complex, cross-domain tasks.

-  [AI Competitor Intelligence Agent Team](#) - Structured competitor teardowns built from their own websites
-  [AI Finance Agent Team](#) - A financial analyst team in 20 lines of Python
-  [AI Game Design Agent Team](#) - Full game concepts from a swarm of design specialists
-  [AG2 Adaptive Research Team](#) - Agent teamwork with routing and fallback, built on AG2
-  [AI Legal Agent Team \(Cloud & Local\)](#) - Research, contract analysis, and strategy from a full legal bench
-  [AI Recruitment Agent Team](#) - Resume screening to interview scheduling, end to end
-  [AI Real Estate Agent Team](#) - Property search, market analysis, and recommendations
-  [AI Services Agency \(CrewAI\)](#) - A digital agency that scopes and plans your software project
-  [AI Teaching Agent Team](#) - A faculty of agents that builds your complete learning path
-  [Multimodal Coding Agent Team](#) - Snap a photo of a coding problem, get a sandboxed solution
-  [Multimodal Design Agent Team](#) - Design critiques from a Gemini-powered expert panel
-  [Multimodal UI/UX Feedback Agent Team](#) - Landing page feedback plus an auto-generated improved version
-  [AI Travel Planner Agent Team](#) - A complete trip itinerary, crafted by a team

Voice AI Agents

Speech-in, speech-out agents using real-time voice APIs.

-  [AI Audio Tour Agent](#) - Self-guided audio tours from your location, interests, and pace
-  [Customer Support Voice Agent](#) - Voice answers grounded in your own docs
-  [Insurance Claim Live Agent Team](#) - Real-time voice claim intake with Gemini Live
-  [Voice RAG Agent \(OpenAI SDK\)](#) - Ask your PDFs questions, hear the answers
-  [OpenSource Voice Dictation Agent \(Wispr Flow clone\)](#) external - Open-source dictation that types where you talk

Generative UI and Agentic Frontends

Agents that render interactive UI components, not just text: forms, cards, charts, editable plans.

-  [Generative UI Starter Project](#) - A chat-driven kanban board you and the agent work together
-  [AI Financial Coach Agent](#) - Budget, savings, and debt plans rendered as interactive cards
-  [AI Dashboard Canvas Agent](#) - Describe a dashboard in chat, charts assemble on a live canvas

-  [AI MCP App Builder](#) - Describe an MCP app, get a live sandboxed instance back
-  [MCP Apps Generative UI Showcase](#) - MCP apps that render real interactive UI, flight search included
-  [AI Shadcn Component Generator](#) - Chat your way to production-ready shadcn components
-  [AI Deep Research Agent](#) - Research where every tool call renders as a live workspace card

Autonomous Game-Playing Agents

Agents that play games end-to-end: reasoning, strategy, and action.

-  [AI 3D Pygame Agent](#) - DeepSeek R1 writes PyGame code, browser agents run it live
-  [AI Chess Agent](#) - Agent White vs Agent Black with validated moves
-  [AI Tic-Tac-Toe Agent](#) - Two different LLMs battle it out, move by move

MCP AI Agents

Agents that connect to external tools and data via Model Context Protocol.

-  [Browser MCP Agent](#) - Drive a real browser with natural language over MCP
-  [GitHub MCP Agent](#) - Explore and analyze any repo in plain English
-  [Notion MCP Agent](#) - Talk to your Notion pages from the terminal
-  [AI Travel Planner MCP Agent](#) - Itineraries built on live Airbnb and Google Maps data
-  [Multi-MCP Agent Router](#) - Specialist agents, each wired to its own MCP server
-  [OpenAI Remote MCP Tool Bridge](#) - Connect OpenAI function calling directly to a remote MCP server

RAG (Retrieval Augmented Generation)

Retrieval pipelines, from simple chains to agentic and multi-source.

-  [Agentic RAG with Embedding Gemma](#) - Fully local agentic RAG with EmbeddingGemma and Llama 3.2
-  [Agentic RAG with Reasoning](#) - Watch the agent's step-by-step reasoning as it retrieves
-  [AI Blog Search \(RAG\)](#) - Agentic search over blog content, built on LangGraph
-  [Autonomous RAG](#) - GPT-4o answers from your PDFs, falls back to web search
-  [Contextual AI RAG Agent](#) - Managed RAG: datastore to grounded chat in minutes
-  [Corrective RAG \(CRAG\)](#) - Retrieval that grades itself and retries before answering
-  [Typed Agentic RAG with Pydantic AI](#) - Validated answers with exact citations, or a refusal when evidence is weak
-  [Deepseek Local RAG Agent](#) - Local DeepSeek reasoning over your own documents

- 🤖 [Gemini Agentic RAG](#) - Query rewriting and web fallback with Gemini Flash Thinking
- 🗄️ [Hybrid Search RAG \(Cloud\)](#) - Keyword plus vector search feeding Claude
- 🔄 [Llama 3.1 Local RAG](#) - Chat with any webpage, fully offline
- 💻 [Local Hybrid Search RAG](#) - Hybrid search with everything running on your machine
- 🧬 [Multimodal Agentic RAG](#) - Text, PDFs, images, audio, and video, answered with citations
- 🐘 [Local RAG Agent](#) - Llama 3.2 and Qdrant, no API keys required
- 🧩 [RAG-as-a-Service](#) - A production RAG service in under 50 lines
- ⚡ [RAG Agent with Cohere](#) - Command R7B retrieval with web-search fallback
- 🧬 [Basic RAG Chain](#) - The minimal retrieval pipeline, applied to pharma research
- 📊 [RAG with Database Routing](#) - Routes each question to the right database automatically
- 🖼️ [Vision RAG](#) - Ask questions about images and PDF pages with Embed-4
- 🩺 [RAG Failure Diagnostics Clinic](#) - Find out why your RAG pipeline is wrong, systematically
- 📄 [Knowledge Graph RAG with Citations](#) - Multi-hop answers with verifiable source attribution

💾 LLM Apps with Memory

Agents and chatbots that remember conversations and user state across sessions.

- 💾 [AI ArXiv Agent with Memory](#) - Paper search that remembers your research interests
- ✈️ [AI Travel Agent with Memory](#) - A travel assistant that remembers your preferences
- 💬 [Llama3 Stateful Chat](#) - Session-persistent chat with Llama 3
- 📝 [LLM App with Personalized Memory](#) - A chatbot that keeps context across conversations
- 📱 [Local ChatGPT Clone with Memory](#) - Fully local, with a personal memory per user
- 🗣️ [Multi-LLM Application with Shared Memory](#) - Different models, one shared conversation memory

💬 Chat with X

Turn any data source into a chat interface.

- 💬 [Chat with GitHub \(GPT & Llama3\)](#) - Any repo, answered in 30 lines of RAG
- 📧 [Chat with Gmail](#) - Ask your inbox questions
- 📄 [Chat with PDF \(GPT & Llama3\)](#) - The classic, in 30 lines of Python
- 📖 [Chat with Research Papers \(ArXiv\) \(GPT & Llama3\)](#) - Explore arXiv conversationally with GPT-4o
- 📰 [Chat with Substack](#) - Chat with any newsletter's archive
- 🎥 [Chat with YouTube Videos](#) - Ask videos questions via their transcripts

LLM Optimization Tools

Reduce token usage, context size, and API cost without losing quality.

-  [Toonify Token Optimization](#) - Reduce LLM API costs by 30-60% using TOON format
-  [Headroom Context Optimization](#) - Reduce LLM API costs by 50-90%

LLM Fine-tuning

End-to-end fine-tuning recipes for open-source models.

-  [Gemma 3 Fine-tuning](#) - 4-bit LoRA with Unsloth, small and readable
-  [Llama 3.2 Fine-tuning](#) - Fine-tune in 30 lines, free on Colab

AI Agent Framework Crash Courses

Deep-dive tutorials on the major agent frameworks.

- [Google ADK Crash Course](#) - Starter agent, structured outputs, tools (built-in, function, third-party, MCP), memory, callbacks, plugins, and multi-agent patterns. Model-agnostic.
- [OpenAI Agents SDK Crash Course](#) - Starter agent, function calling, structured outputs, tools, memory, evaluation, handoffs, swarm orchestration, and routing logic.

Getting Started with uv

This project now uses [uv](#) for fast, reliable Python package management.

Prerequisites

1. **Install uv** (if not already installed)

```
# macOS and Linux
curl -LsSf https://astral.sh/uv/install.sh | sh

# Windows
powershell -c "irm https://astral.sh/uv/install.ps1 | iex"

# Or with pip
pip install uv
```

Quick Start

1. Clone the repository

```
git clone https://github.com/Shubhamsaboo/awesome-llm-apps.git
cd awesome-llm-apps
```

2. Install all dependencies with uv

```
# Install core dependencies only
uv sync

# Or install with all optional dependencies
uv sync --all-extras

# Or install specific extras (agents, rag, web, mcp, etc.)
uv sync --extra agents --extra rag
```

3. Navigate to the desired project directory

```
cd starter_ai_agents/ai_travel_agent
```

4. Run the project using uv

```
# Run Python scripts
uv run python script_name.py

# Run Streamlit apps
uv run streamlit run app.py
```

Alternative: Traditional pip Installation

If you prefer using pip, you can still install dependencies for individual projects:

```
cd awesome-llm-apps/starter_ai_agents/ai_travel_agent
pip install -r requirements.txt
```

Why uv?

- ⚡ **10-100x faster** than pip for dependency resolution and installation
- 🔒 **Reliable** with a lockfile (`uv.lock`) ensuring reproducible environments

-  **Simple** single command to set up the entire project
-  **Compatible** works alongside existing pip/requirements.txt workflows

Available Dependency Groups

The project includes optional dependency groups you can install as needed:

- `agents` - AI Agent frameworks (agno, openai-agents, mcp-agent)
- `rag` - RAG and vector stores (lancedb, pypdf, faiss)
- `web` - Web scraping and API tools
- `mcp` - MCP agent tools
- `local-llm` - Local LLM support (ollama)
- `dev` - Development tools (pytest, ruff, mypy)
- `all` - All dependencies

Follow the project-specific instructions in each project's `README.md` file to set up and run the app.

★ **Star the repo** to get notified when new templates drop.

[Deutsch](#) ·
[Español](#) ·
[français](#) ·
[日本語](#) ·
[한국어](#) ·
[Português](#) ·
[Русский](#) ·
[中文](#)

Apache-2.0 · See [LICENSE](#) · Fork it, ship it, sell it.